

Stock prediction, linear regression/SVM

Jacob Mengistu

2024-09-23

Load Libraries

```
library(ggplot2)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(ggpubr)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v forcats   1.0.0      v stringr   1.5.1
## v lubridate 1.9.3      v tibble   3.2.1
## v purrr     1.0.2      v tidyr    1.3.1
## v readr     2.1.5
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift
```

```
library(broom)
library(AICcmodavg)
library(plyr)
```

```
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
## -----
##
## Attaching package: 'plyr'
##
## The following object is masked from 'package:purrr':
##
##   compact
##
## The following object is masked from 'package:ggpubr':
##
##   mutate
##
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
library(rstatix)
```

```
##
## Attaching package: 'rstatix'
##
## The following objects are masked from 'package:plyr':
##
##   desc, mutate
##
## The following object is masked from 'package:stats':
##
##   filter
```

```
library(stats)
library("lubridate")
library(e1071)
```

Load dataset and view dataframe

```
df<-read.csv('ADANIPORTS.csv')
head(df)
```

```
##      Date      Symbol Series Prev.Close  Open   High Low Last   Close
## 1 2007-11-27 MUNDRAPORT    EQ    440.00 770.00 1050.00 770  959  962.90
```

```
## 2 2007-11-28 MUNDRAPORT      EQ      962.90 984.00 990.00 874 885 893.90
## 3 2007-11-29 MUNDRAPORT      EQ      893.90 909.00 914.75 841 887 884.20
## 4 2007-11-30 MUNDRAPORT      EQ      884.20 890.00 958.00 890 929 921.55
## 5 2007-12-03 MUNDRAPORT      EQ      921.55 939.75 995.00 922 980 969.30
## 6 2007-12-04 MUNDRAPORT      EQ      969.30 985.00 1056.00 976 1049 1041.45
##      VWAP      Volume      Turnover Trades Deliverable.Volume X.Deliverble
## 1  984.72 27294366 2.687719e+15      NA      9859619      0.3612
## 2  941.38 4581338 4.312765e+14      NA      1453278      0.3172
## 3  888.09 5124121 4.550658e+14      NA      1069678      0.2088
## 4  929.17 4609762 4.283257e+14      NA      1260913      0.2735
## 5  965.65 2977470 2.875200e+14      NA      816123      0.2741
## 6 1015.39 4849250 4.923867e+14      NA      1537667      0.3171
```

Create a year column to group by

```
df$year <- year(df$Date)
```

Check and remove rows w/null values

```
sum(is.na(df))
```

```
## [1] 866
```

```
df <- na.omit(df)
head(df)
```

```
##      Date      Symbol Series Prev.Close  Open  High  Low  Last  Close
## 867 2011-06-01 MUNDRAPORT      EQ      161.45 162.10 165.70 161.25 163.50 164.00
## 868 2011-06-02 MUNDRAPORT      EQ      164.00 164.00 165.15 160.15 161.15 161.25
## 869 2011-06-03 MUNDRAPORT      EQ      161.25 161.50 162.80 159.20 161.00 161.05
## 870 2011-06-06 MUNDRAPORT      EQ      161.05 160.50 161.10 159.05 160.00 159.85
## 871 2011-06-07 MUNDRAPORT      EQ      159.85 159.85 162.75 156.35 157.00 157.25
## 872 2011-06-08 MUNDRAPORT      EQ      157.25 157.40 158.70 153.30 155.65 154.90
##      VWAP      Volume      Turnover Trades Deliverable.Volume X.Deliverble year
## 867 164.08 2574106 4.223703e+13 19171      1271255      0.4939 2011
## 868 162.17 1699298 2.755678e+13 16176      791462      0.4658 2011
## 869 161.02 1185817 1.909361e+13 14810      722154      0.6090 2011
## 870 160.09 546378 8.746905e+12 7071      386144      0.7067 2011
## 871 158.52 2193466 3.477027e+13 17865      1425849      0.6500 2011
## 872 155.48 1431303 2.225382e+13 15027      931903      0.6511 2011
```

Summarizing open, close, high, low by year

Open summary by year

```
Open_summary <- df %>% group_by(year) %>% summarize(mean = mean(Open), median = median(Open), min= min
print(Open_summary)
```

```
##           mean  median min max
## 1 284.9406 299.475 108 857
```

Close summary by year


```
Close_summary <- df %>% group_by(year) %>% summarize(mean = mean(Close), median = median(Close), min = min(Close), max = max(Close))
print(Close_summary)
```

```
##           mean median min    max
## 1 284.6273 298.75 108 835.55
```

Random plots for fun

compare opening price through the years w line graph

```
ggplot(data = df, aes(x = year, y = Open)) + geom_line() + scale_x_continuous(breaks = seq(min(df$year), max(df$year), by = 1))
```




Create a scatter plot for open price

```
ggplot(data = df, aes(x = year, y = Open, color = year)) + geom_point() + scale_x_continuous(breaks = seq(min(df$year), max(df$year), by = 1))
```


Scatterplot of year and close price

```
ggplot(df, aes(x = year, y = Close)) + geom_point()
```



Histogram for close price

```
ggplot(df, aes(x = Close)) + geom_histogram(binwidth = 1) #Line plot of open prices plot(df$Open, type='l')
```



Create a target column where 'Up' if $\text{Close} > \text{Open}$, 'Down' otherwise for SVM

```
df$target <- ifelse(df$Close > df$Open, "Up", "Down")
```

Convert the target column to a factor

```
df$target <- as.factor(df$target)
```

Check the levels of the factor to make sure they are "Up" and "Down"

```
levels(df$target)
```

```
## [1] "Down" "Up"
```

Split the data into training and test sets

Set seed

```
set.seed(123)
```

Training data

```
train_indices <- sample(1:nrow(df), 0.8*nrow(df))
train <- df[train_indices, ]
```

Testing data


```
test <- df[-train_indices, ]
```


Fit the linear regression model for close based on previous close prev.close


```
model_fit <- lm(Close ~ Prev.Close + Open, data = train)
close_predictions <- predict(model_fit, newdata = test)
summary(model_fit)
```


```
##
## Call:
## lm(formula = Close ~ Prev.Close + Open, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -65.674  -3.171  -0.259   2.974  83.275
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.54987    0.39354   1.397  0.16250
## Prev.Close    0.15665    0.05183   3.023  0.00254 **
## Open          0.84054    0.05166  16.271 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.929 on 1961 degrees of freedom
## Multiple R-squared:  0.9969, Adjusted R-squared:  0.9969
## F-statistic: 3.108e+05 on 2 and 1961 DF, p-value: < 2.2e-16
```

```
plot(model_fit)
```









Evaluate the model using RMSE and R-squared

```
mse <- mean((test$Close - close_predictions)^2)

rmse <- sqrt(mse)

mae <- mean(abs(test$Close - close_predictions))

r_squared <- cor(test$Close, close_predictions)^2

mse <- mean((test$Close - close_predictions)^2)

rmse <- sqrt(mse)

mae <- mean(abs(test$Close - close_predictions))

r_squared <- cor(test$Close, close_predictions)^2
```

Print results

```
print(paste("MSE for closing price:", mse))
```

```
## [1] "MSE for closing price: 40.9712553764558"
```

```
print(paste("RMSE for closing price:", rmse))
```

```
## [1] "RMSE for closing price: 6.40087926588651"
```

```
print(paste("MAE for closing price:", mae))
```

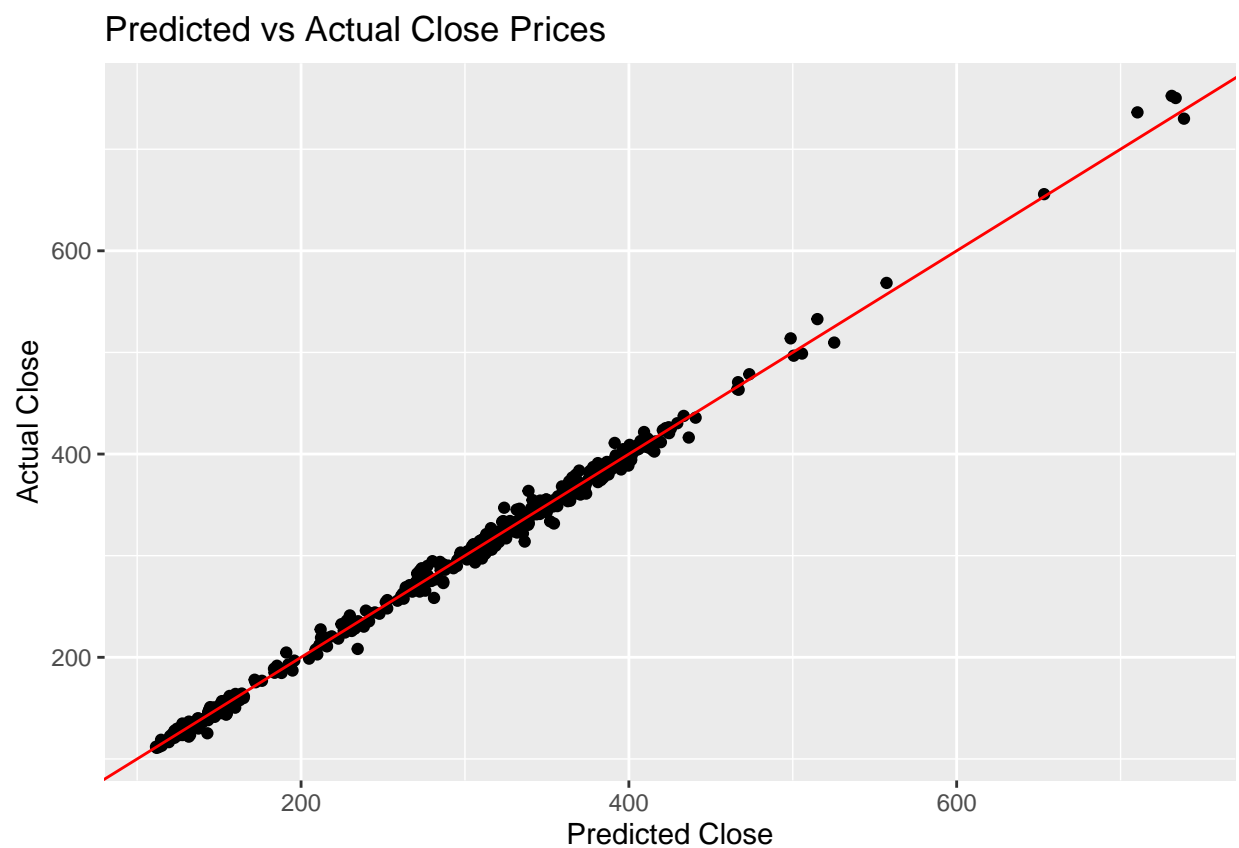
```
## [1] "MAE for closing price: 4.60259732369209"
```

```
print(paste("R-squared for closing price:", r_squared))
```

```
## [1] "R-squared for closing price: 0.996854728114382"
```

Plot actual vs predicted close prices

```
ggplot(data = test, aes(x = close_predictions, y = Close)) + geom_point() + geom_abline(slope = 1, intercept = 0)
```



Train the SVM model using the ‘open’ and ‘close’ variables as predictors to see if it can predict price going up or down

*Not a very good model, could keep tweaking but would rather try others

*Tried to optimize, not very successful

```
train_control <- trainControl(method = "cv", number = 10)
tune_grid <- expand.grid(C = seq(0.01, 1, by = 0.1))
```

*Cross-validation before it was guessing at .5, randomly, .5 after

SVM Model

```
svm_model <- train(target ~ Open + Prev.Close, data = train, method = "svmLinear", trControl = train_c
```

Make predictions on the test set using the trained model

```
predictions <- predict(svm_model, newdata = test)
```

Evaluate the model’s performance using a confusion matrix

```
confusion_matrix <- confusionMatrix(predictions, test$target)
print(confusion_matrix)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Down  Up
##      Down  246  246
##      Up     0   0
##
##              Accuracy : 0.5
##              95% CI : (0.4549, 0.5451)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 0.518
##
##              Kappa : 0
##
##      Mcnemar's Test P-Value : <2e-16
```

```
##
##      Sensitivity : 1.0
##      Specificity : 0.0
##      Pos Pred Value : 0.5
##      Neg Pred Value : NaN
##      Prevalence : 0.5
##      Detection Rate : 0.5
##      Detection Prevalence : 1.0
##      Balanced Accuracy : 0.5
##
##      'Positive' Class : Down
##
```